

# Financial Data Analysis with Python

Instructor: Luping Yu

Mar 8, 2022

---


## Lecture 03. Data Loading and Cleaning

Accessing data is a necessary first step for using most of the tools in this course. I'm going to be focused on data input and output using pandas.

---

### Reading and writing data in text format

pandas features a number of functions for reading **tabular data** as a **DataFrame** object.



The following table summarizes some of them, though **read\_csv** and **read\_table** are likely the ones you'll use the most.

Function	Description
<b>read_csv</b>	Load delimited data from a file, URL, or file-like object; use comma as default delimiter
<b>read_excel</b>	Read tabular data from an Excel XLS or XLSX file
<b>read_stata</b>	Read a dataset from Stata file format
<b>read_sas</b>	Read a SAS dataset stored in one of the SAS system's custom storage formats
<b>read_html</b>	Read all tables found in the given HTML document
<b>read_json</b>	Read data from a JSON (JavaScript Object Notation) string representation
<b>read_pickle</b>	Read an arbitrary object stored in Python pickle format
<b>read_sql</b>	Read the results of a SQL query (using SQLAlchemy) as a pandas DataFrame

---

### Reading and Writing .csv (comma-separated values)

.csv is a delimited text file that uses a **comma** to separate values. A CSV file typically stores **tabular data** (numbers and text) in **plain text**.

Let's start with a small comma-separated (CSV) text file: [ex1.csv](#)

```
In [244... cat examples/ex1.csv
```

```
a,b,c,d,message
1,2,3,4,hello
5,6,7,8,world
9,10,11,12,foo
```

Here I used the Linux/macOS **cat** shell command to print the raw contents of the file to the screen. If you're on Windows, you can use **type** instead of cat to achieve the same effect.

```
In [245... import pandas as pd

pd.read_csv('examples/ex1.csv') # 相对路径
```

```
Out[245]:
```

	a	b	c	d	message
0	1	2	3	4	hello
1	5	6	7	8	world
2	9	10	11	12	foo

```
In [246... # 绝对路径 (注意windows和mac绝对路径命名方式不同)
pd.read_csv('/Users/luping/desktop/teaching/examples/ex1.csv')
```

```
Out[246]:
```

	a	b	c	d	message
0	1	2	3	4	hello
1	5	6	7	8	world
2	9	10	11	12	foo

pandas.read\_csv perform type inference. That means you don't necessarily have to specify which columns are numeric, integer, boolean, or string:

```
In [247... df = pd.read_csv('/Users/luping/desktop/teaching/examples/ex1.csv')

df.dtypes
```

```
Out[247]:
```

a	int64
b	int64
c	int64
d	int64
message	object
dtype:	object

A file will not always have a header row. Consider this file: [ex2.csv](#)

```
In [248... cat examples/ex2.csv
```

```
1,2,3,4,hello
5,6,7,8,world
9,10,11,12,foo
```

To read this file, you have a couple of options. You can allow pandas to assign default column names, or you can specify names yourself:

```
In [249... pd.read_csv('examples/ex2.csv', header=None)
```

```
Out[249]:
```

	0	1	2	3	4
0	1	2	3	4	hello
1	5	6	7	8	world
2	9	10	11	12	foo

```
In [250... pd.read_csv('examples/ex2.csv', names=['a', 'b', 'c', 'd', 'message'])
```

```
Out[250]:
```

	a	b	c	d	message
0	1	2	3	4	hello
1	5	6	7	8	world
2	9	10	11	12	foo

Suppose you wanted the message column to be the index of the returned DataFrame. You can either indicate you want the column at index 4 or named 'message' using the `index_col` argument:

```
In [251... names = ['a', 'b', 'c', 'd', 'message']

pd.read_csv('examples/ex2.csv', names=names, index_col='message')
```

```
Out[251]:
```

message	a	b	c	d
hello	1	2	3	4
world	5	6	7	8
foo	9	10	11	12

The parser functions have many additional arguments to help you handle the wide variety of exception file formats that occur. Consider this file: [ex3.csv](#)

```
In [252... cat examples/ex3.csv

# Hey!
a,b,c,d,message
# Author: Luping Yu
# 厦门大学管理学院
1,2,3,4,hello
5,6,7,8,world
9,10,11,12,foo
```

For example, you can skip the first, third, and fourth rows of a file with **skiprows**:

```
In [253... pd.read_csv('examples/ex3.csv', skiprows=[0, 2, 3])
```

```
Out[253]:
```

	a	b	c	d	message
0	1	2	3	4	hello
1	5	6	7	8	world
2	9	10	11	12	foo

Handling missing values is an important and frequently nuanced part of the file parsing process. Missing data is usually either not present (empty string) or marked by some sentinel value. Consider this file: [ex4.csv](#)

```
In [254... cat examples/ex4.csv

something,a,b,c,d,message
one,1,2,3,4,NA
two,5,6,,8,world
three,9,10,11,12,foo
```

By default, pandas uses a set of commonly occurring sentinels, such as **NA** and **NULL**:

```
In [255... pd.read_csv('examples/ex4.csv')
```

```
Out[255]:
```

	something	a	b	c	d	message
0	one	1	2	3.0	4	NaN
1	two	5	6	NaN	8	world
2	three	9	10	11.0	12	foo

```
In [256... df = pd.read_csv('examples/ex4.csv')

pd.isnull(df)
```

```
Out[256]:
```

	something	a	b	c	d	message
0	False	False	False	False	False	True
1	False	False	False	True	False	False
2	False	False	False	False	False	False

Data can also be exported to a delimited format. Let's consider one of the CSV files read before:

```
In [257... df
```

```
Out[257]:
```

	something	a	b	c	d	message
0	one	1	2	3.0	4	NaN
1	two	5	6	NaN	8	world
2	three	9	10	11.0	12	foo

Using DataFrame's `to_csv` method, we can write the data out to a comma-separated file:

```
In [258... df.to_csv('examples/out1.csv')
```

```
In [259... cat examples/out1.csv  
  
,something,a,b,c,d,message  
0,one,1,2,3.0,4,  
1,two,5,6,,8,world  
2,three,9,10,11.0,12,foo
```

With no other options specified, both the row and column labels are written. Both of these can be disabled:

```
In [260... df.to_csv('examples/out2.csv', index=False, header=False)
```

```
In [261... cat examples/out2.csv  
  
one,1,2,3.0,4,  
two,5,6,,8,world  
three,9,10,11.0,12,foo
```

---

## Parameters of data loading functions

The optional arguments for these functions may fall into a few categories:

- Indexing

Can treat one or more columns as the returned DataFrame, and whether to get column names from the file, the user, or not at all.

- Type inference and data conversion

This includes the user-defined value conversions and custom list of missing value markers.

- Datetime parsing

Includes combining capability, including combining date and time information spread over multiple columns into a single column in the result.

- Iterating

Support for iterating over chunks of very large files.

- Unclean data issues

Skipping rows or a footer, comments, or other minor things like numeric data with thousands separated by commas.

Because of how messy data in the real world can be, some of the data loading functions (especially `read_csv`) have grown very complex in their options over time. It's normal to feel overwhelmed by the number of different parameters (`read_csv` has over 50 as of this writing). The **online pandas documentation** has many examples about how each of them works, so if you're struggling to read a particular file, there might be a similar enough example to help you find the right parameters.

API reference (pandas documentation) of `read_csv`:

[https://pandas.pydata.org/docs/reference/api/pandas.read\\_csv.html](https://pandas.pydata.org/docs/reference/api/pandas.read_csv.html)

---

## Reading microsoft excel files

pandas also supports reading tabular data stored in Excel 2003 (and higher) files using either the `ExcelFile` class or `pandas.read_excel` function. Internally these tools use the add-on packages **xlrd** and **openpyxl** to read XLS and XLSX files, respectively. You may need to install these manually with pip or conda.

To use `ExcelFile`, pass the filename to `pandas.read_excel`:

```
In [262]: df = pd.read_excel('examples/ex1.xlsx', 'Sheet1')
df
```

```
Out[262]:
```

	a	b	c	d	message
0	1	2	3	4	hello
1	5	6	7	8	world
2	9	10	11	12	foo

To write pandas data to Excel format, you can pass a file path to **to\_excel**:

```
In [263]: df.to_excel('examples/out1.xlsx')
```

---

## Data cleaning and preparation

During the course of doing data analysis and modeling, a significant amount of time is spent on data preparation: loading, cleaning, transforming, and rearranging. Such tasks are often reported to take up 80% or more of an analyst's time. Sometimes the way that data is stored in files or databases is not in the right format for a particular task.

Fortunately, pandas, along with the built-in Python language features, provides you with a high-level, flexible, and fast set of tools to enable you to manipulate data into the right form.

Much of the design and implementation of pandas has been driven by the needs of real-world applications.

## Handling Missing Data

Missing data occurs commonly in many data analysis applications.

For numeric data, pandas uses the floating-point value **NaN** (Not a Number) to represent missing data.

In pandas, we've adopted a convention used in the R programming language by referring to missing data as **NA**, which stands for **not available**. When cleaning up data for analysis, it is often important to do analysis on the missing data itself to identify data collection problems or potential biases in the data caused by missing data.

With DataFrame objects, you may want to drop rows or columns that are all NA or only those containing any NAs.

`dropna` by default drops any row containing a missing value:

```
In [264... df = pd.DataFrame([[1., 6.5, 3.],
                    [1., None, None],
                    [None, None, None],
                    [None, 6.5, 3.]])

df
```

```
Out[264]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
2	NaN	NaN	NaN
3	NaN	6.5	3.0

```
In [265... df.dropna()
```

```
Out[265]:
```

	0	1	2
0	1.0	6.5	3.0

Passing `how='all'` will only drop rows that are all NA:

```
In [266... df.dropna(how='all')
```

```
Out[266]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
3	NaN	6.5	3.0

To drop columns in the same way, pass `axis=1`:

```
In [267... df[3] = None

df
```

```
Out[267]:
```

	0	1	2	3
0	1.0	6.5	3.0	None
1	1.0	NaN	NaN	None
2	NaN	NaN	NaN	None
3	NaN	6.5	3.0	None

```
In [268... df.dropna(axis=1, how='all')
```

```
Out[268]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
2	NaN	NaN	NaN
3	NaN	6.5	3.0

## Filling In Missing Data

Rather than filtering out missing data (and potentially discarding other data along with it), you may want to fill in the “holes” in any number of ways. For most purposes, the `fillna` method is the workhorse function to use.

Calling `fillna` with a constant replaces missing values with that value:

```
In [269]: df = pd.DataFrame([[10, 10, 20, 60],
                             [8, 5, 12, 48],
                             [6, 10, 14, None],
                             [None, None, None, None],
                             [None, None, 10, 30]],
                             columns=['平时分', '小作业', '大作业', '期末'])

df
```

```
Out[269]:
```

	平时分	小作业	大作业	期末
0	10.0	10.0	20.0	60.0
1	8.0	5.0	12.0	48.0
2	6.0	10.0	14.0	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	10.0	30.0

```
In [270]: df.fillna(0)
```

```
Out[270]:
```

	平时分	小作业	大作业	期末
0	10.0	10.0	20.0	60.0
1	8.0	5.0	12.0	48.0
2	6.0	10.0	14.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	10.0	30.0

Calling `fillna` with a dict, you can use a different fill value for each column:

```
In [271]: df
```



```
Out[271]:
```

	平时分	小作业	大作业	期末
0	10.0	10.0	20.0	60.0
1	8.0	5.0	12.0	48.0
2	6.0	10.0	14.0	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	10.0	30.0

```
In [272]: df.fillna({'平时分': 5, '期末': 30})
```

```
Out[272]:
```

	平时分	小作业	大作业	期末
0	10.0	10.0	20.0	60.0
1	8.0	5.0	12.0	48.0
2	6.0	10.0	14.0	30.0
3	5.0	NaN	NaN	30.0
4	5.0	NaN	10.0	30.0

The interpolation methods can be used with fillna:

```
In [273]: df
```

```
Out[273]:
```

	平时分	小作业	大作业	期末
0	10.0	10.0	20.0	60.0
1	8.0	5.0	12.0	48.0
2	6.0	10.0	14.0	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	10.0	30.0

```
In [274]: df.fillna(method='ffill')
```

```
Out[274]:
```

	平时分	小作业	大作业	期末
0	10.0	10.0	20.0	60.0
1	8.0	5.0	12.0	48.0
2	6.0	10.0	14.0	48.0
3	6.0	10.0	14.0	48.0
4	6.0	10.0	10.0	30.0

With fillna you can do lots of other things with a little creativity. For example, you might pass the mean or median value of a Series:

```
In [275]: df
```

```
Out[275]:
```

	平时分	小作业	大作业	期末
0	10.0	10.0	20.0	60.0
1	8.0	5.0	12.0	48.0
2	6.0	10.0	14.0	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	10.0	30.0

```
In [276]: df.fillna(df.mean())
```

```
Out[276]:
```

	平时分	小作业	大作业	期末
0	10.0	10.000000	20.0	60.0
1	8.0	5.000000	12.0	48.0
2	6.0	10.000000	14.0	46.0
3	8.0	8.333333	14.0	46.0
4	8.0	8.333333	10.0	30.0

## Removing Duplicates

Duplicate rows may be found in a DataFrame for any number of reasons. Here is an example:

```
In [277]: df = pd.DataFrame({'k1': ['one', 'two'] * 3 + ['two'],
                             'k2': [1, 1, 2, 3, 3, 4, 4]})

df
```

```
Out[277]:
```

	k1	k2
0	one	1
1	two	1
2	one	2
3	two	3
4	one	3
5	two	4
6	two	4

The DataFrame method `df.duplicated()` returns a boolean Series indicating whether each row is a duplicate (has been observed in a previous row) or not:

```
In [278]: df.duplicated()
```

```
Out[278]: 0    False
          1    False
          2    False
          3    False
          4    False
          5    False
          6     True
          dtype: bool
```

Relatedly, `drop_duplicates` returns a DataFrame where the duplicated array is False:

```
In [279... df.drop_duplicates()
```

```
Out[279]:
```

	k1	k2
0	one	1
1	two	1
2	one	2
3	two	3
4	one	3
5	two	4

Both of these methods by default consider all of the columns; alternatively, you can specify any subset of them to detect duplicates. Suppose we had an additional column of values and wanted to filter duplicates only based on the 'k1' column:

```
In [280... df['v1'] = range(7)
df
```

```
Out[280]:
```

	k1	k2	v1
0	one	1	0
1	two	1	1
2	one	2	2
3	two	3	3
4	one	3	4
5	two	4	5
6	two	4	6

```
In [281... df.drop_duplicates(['k1'])
```

```
Out[281]:
```

	k1	k2	v1
0	one	1	0
1	two	1	1

`drop_duplicates` by default keep the first observed value combination. Passing `keep='last'` will return the last one:

```
In [282... df
```

```
Out[282]:
```

	k1	k2	v1
0	one	1	0
1	two	1	1
2	one	2	2
3	two	3	3
4	one	3	4
5	two	4	5
6	two	4	6

```
In [283]: df.drop_duplicates(['k1', 'k2'], keep='last')
```

```
Out[283]:
```

	k1	k2	v1
0	one	1	0
1	two	1	1
2	one	2	2
3	two	3	3
4	one	3	4
6	two	4	6

---

## Replacing Values

Filling in missing data with the `fillna` method is a special case of more general value replacement. `replace` provides a simpler and more flexible way to do so. Let's consider this Series:

```
In [284]: df = pd.Series([3321, 1742.5, -999, 381.2, 13.42])
df
```

```
Out[284]:
```

0	3321.00
1	1742.50
2	-999.00
3	381.20
4	13.42

dtype: float64

The -999 values might be sentinel values for missing data. To replace these with NA values that pandas understands, we can use `replace`, producing a new Series:

```
In [285]: df.replace(-999, None)
```

```
Out[285]:
```

0	3321.0
1	1742.5
2	None
3	381.2
4	13.42

dtype: object

---

## Vectorized string functions in pandas

Cleaning up a messy dataset for analysis often requires a lot of string munging and regularization. To complicate matters, a column containing strings will sometimes have missing data:

```
In [286]: df = pd.Series({'Dave': 'dave@google.com',  
                        'Jack': 'jack@xmu.edu.cn',  
                        'Steve': 'steve@gmail.com',  
                        'Rose': 'rose@xmu.edu.cn',  
                        'Tony': None})  
  
df
```

```
Out[286]: Dave      dave@google.com  
Jack      jack@xmu.edu.cn  
Steve     steve@gmail.com  
Rose      rose@xmu.edu.cn  
Tony      None  
dtype: object
```

To cope with this, Series has array-oriented methods for string operations that skip NA values. These are accessed through Series's `str` attribute; for example, we could check whether each email address has 'xmu.edu' in it with [`str.contains`](#):

```
In [287]: df.str.contains('xmu.edu')
```

```
Out[287]: Dave      False  
Jack      True  
Steve     False  
Rose      True  
Tony      None  
dtype: object
```

You can similarly slice strings using this syntax:

```
In [288]: df.str[:5]
```

```
Out[288]: Dave      dave@  
Jack      jack@  
Steve     steve  
Rose      rose@  
Tony      None  
dtype: object
```

```
In [289]: df.str.split('@')
```

```
Out[289]: Dave      [dave, google.com]  
Jack      [jack, xmu.edu.cn]  
Steve     [steve, gmail.com]  
Rose      [rose, xmu.edu.cn]  
Tony      None  
dtype: object
```

```
In [290]: df.str.split('@').str.get(0)
```

```
Out[290]: Dave      dave  
Jack      jack  
Steve     steve  
Rose      rose  
Tony      None  
dtype: object
```

Partial listing of vectorized string methods.

Method	Description
cat	Concatenate strings element-wise with optional delimiter
contains	Return boolean array if each string contains pattern/regex
count	Count occurrences of pattern
extract	Use a regular expression with groups to extract one or more strings from a Series of strings
endswith	Equivalent to x.endswith(pattern) for each element
startswith	Equivalent to x.startswith(pattern) for each element
findall	Compute list of all occurrences of pattern/regex for each string
get	Index into each element (retrieve i-th element)
join	Join strings in each element of the Series with passed separator
len	Compute length of each string
lower,upper	Convert cases;equivalent to x.lower() or x.upper() for each element
match	Use re.match with the passed regular expression on each element
replace	Replace occurrences of pattern/regex with some other string
slice	Slice each string in the Series
split	Split strings on delimiter or regular expression
strip	Trim whitespace from both sides, including newlines