

Financial Data Analysis with Python

Instructor: Luping Yu

Mar 22, 2022

Project 1

It's time to test your new pandas skills! Use the provided .csv file to complete the tasks below!

Student Name 1:

Student Name 2:

Task 1. Bank M&A

Note: failed banks (Bank Name) are bought by another bank (Acquiring Institution).

```
In [138... import pandas as pd
df = pd.read_csv('banklist.csv')
```

```
In [139... df
```

Out[139]:

| | Bank Name | City | ST | CERT | Acquiring Institution | Closing Date | Updated Date |
|-----|---|--------------------|-----|-------|-------------------------------------|--------------|--------------|
| 0 | Fayette County Bank | Saint Elmo | IL | 1802 | United Fidelity Bank, fsb | 26-May-17 | 1-Jun-17 |
| 1 | Guaranty Bank, (d/b/a BestBank in Georgia & Mi... | Milwaukee | WI | 30003 | First-Citizens Bank & Trust Company | 5-May-17 | 1-Jun-17 |
| 2 | First NBC Bank | New Orleans | LA | 58302 | Whitney Bank | 28-Apr-17 | 23-May-17 |
| 3 | Proficio Bank | Cottonwood Heights | UT | 35495 | Cache Valley Bank | 3-Mar-17 | 18-May-17 |
| 4 | Seaway Bank and Trust Company | Chicago | IL | 19328 | State Bank of Texas | 27-Jan-17 | 18-May-17 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 546 | Superior Bank, FSB | Hinsdale | IL | 32646 | Superior Federal, FSB | 27-Jul-01 | 19-Aug-14 |
| 547 | Malta National Bank | Malta | OH | 6629 | North Valley Bank | 3-May-01 | 18-Nov-02 |
| 548 | First Alliance Bank & Trust Co. | Manchester | NH | 34264 | Southern New Hampshire Bank & Trust | 2-Feb-01 | 18-Feb-03 |
| 549 | National State Bank of Metropolis | Metropolis | IL | 3815 | Banterra Bank of Marion | 14-Dec-00 | 17-Mar-05 |
| 550 | Bank of Honolulu | Honolulu | HI | 21029 | Bank of the Orient | 13-Oct-00 | 17-Mar-05 |

551 rows × 7 columns

- Step 1. How many States (ST) are represented in this data set?

```
In [140]: len(df['ST'].drop_duplicates())
```

Out[140]: 44

- Step 2. What are the top 5 states with the most failed banks? (tips: groupby 'ST' and use count)

```
In [141]: df.groupby('ST').count()['Bank Name'].sort_values(ascending=False)[:5]
```

```
Out[141]: ST
GA      93
FL      75
IL      67
CA      41
MN      23
Name: Bank Name, dtype: int64
```

- Step 3. What are the top 5 states with the highest **average** CERT (Certificate) value of failed banks? (tips: groupby 'ST' and use mean)

```
In [142... df.groupby('ST').mean()['CERT'].sort_values(ascending=False)[:5]
```

```
Out[142]: ST
NY      45616.400000
AZ      45201.062500
NJ      44113.428571
KY      43331.000000
NC      39959.000000
Name: CERT, dtype: float64
```

-
- Step 4. What are the top 5 acquiring institutions?

```
In [143... df['Acquiring Institution'].value_counts()[:5]
#df.groupby('Acquiring Institution').count()['Bank Name'].sort_values(ascend
```

```
Out[143]: No Acquirer      31
State Bank and Trust Company    12
First-Citizens Bank & Trust Company  11
Ameris Bank      10
U.S. Bank N.A.      9
Name: Acquiring Institution, dtype: int64
```

-
- Step 5. How many banks has the *State Bank of Texas* acquired?

```
In [144... df[df['Acquiring Institution']=='State Bank of Texas']
```

```
Out[144]:
```

| | Bank Name | City | ST | CERT | Acquiring Institution | Closing Date | Updated Date |
|-----|---------------------------------------|---------|----|-------|-----------------------|--------------|--------------|
| 4 | Seaway Bank and Trust Company | Chicago | IL | 19328 | State Bank of Texas | 27-Jan-17 | 18-May-17 |
| 21 | The National Republic Bank of Chicago | Chicago | IL | 916 | State Bank of Texas | 24-Oct-14 | 6-Jan-16 |
| 450 | Millennium State Bank of Texas | Dallas | TX | 57667 | State Bank of Texas | 2-Jul-09 | 26-Oct-12 |

-
- Step 6. What is the most common city in California for a bank to fail in? (tips: 'ST' == 'CA' and groupby 'City')

```
In [145... df[df['ST']=='CA'].groupby('City')['Bank Name'].count().sort_values(ascendin
```

```
Out[145]: City
Los Angeles      4
Name: Bank Name, dtype: int64
```

-
- Step 7. How many failed banks **don't** have the word "Bank" in their name?

```
In [146... len(df) - sum(df['Bank Name'].str.contains('Bank'))
#sum(df['Bank Name'].str.contains('Bank') == False)
#sum(~df['Bank Name'].str.contains('Bank'))
#sum(df['Bank Name'].apply(lambda name: 'Bank' not in name))
```

Out[146]: 14

-
- Step 8. How many bank names consist of just two words? (e.g. "First Bank" , "Bank Georgia")

```
In [147... sum(df['Bank Name'].str.split(' ').str.len() == 2)
#sum(df['Bank Name'].str.count(' ') + 1 == 2)
```

Out[147]: 113

Task 2. Occupation

```
In [148... import pandas as pd

df = pd.read_csv('occupation.csv', index_col='user_id')
```

```
In [149... df
```

Out[149]:

| | age | gender | occupation | zip_code |
|--|-----|--------|------------|----------|
|--|-----|--------|------------|----------|

| user_id | | | | |
|---------|-----|-----|---------------|-------|
| 1 | 24 | M | technician | 85711 |
| 2 | 53 | F | other | 94043 |
| 3 | 23 | M | writer | 32067 |
| 4 | 24 | M | technician | 43537 |
| 5 | 33 | F | other | 15213 |
| ... | ... | ... | ... | ... |
| 939 | 26 | F | student | 33319 |
| 940 | 32 | M | administrator | 02215 |
| 941 | 20 | M | student | 97229 |
| 942 | 48 | F | librarian | 78209 |
| 943 | 22 | M | student | 77841 |

943 rows × 4 columns

-
- Step 1. What is the mean age per occupation? (tips: groupby 'occupation')

```
In [150... df.groupby('occupation')['age'].mean()
#df.groupby('occupation')['age'].apply('mean')
#df.groupby('occupation')['age'].agg('mean')
```

```
Out[150]: occupation
administrator    38.746835
artist           31.392857
doctor           43.571429
educator         42.010526
engineer         36.388060
entertainment    29.222222
executive        38.718750
healthcare       41.562500
homemaker        32.571429
lawyer           36.750000
librarian        40.000000
marketing        37.615385
none             26.555556
other            34.523810
programmer       33.121212
retired          63.071429
salesman         35.666667
scientist        35.548387
student          22.081633
technician       33.148148
writer           36.311111
Name: age, dtype: float64
```

- Step 2. For each occupation, calculate the minimum and maximum ages.

```
In [151... df.groupby('occupation')['age'].agg(['min', 'max'])
```

Out[151]:

| | min | max |
|---------------|-----|-----|
| occupation | | |
| administrator | 21 | 70 |
| artist | 19 | 48 |
| doctor | 28 | 64 |
| educator | 23 | 63 |
| engineer | 22 | 70 |
| entertainment | 15 | 50 |
| executive | 22 | 69 |
| healthcare | 22 | 62 |
| homemaker | 20 | 50 |
| lawyer | 21 | 53 |
| librarian | 23 | 69 |
| marketing | 24 | 55 |
| none | 11 | 55 |
| other | 13 | 64 |
| programmer | 20 | 63 |
| retired | 51 | 73 |
| salesman | 18 | 66 |
| scientist | 23 | 55 |
| student | 7 | 42 |
| technician | 21 | 55 |
| writer | 18 | 60 |

-
- Step 3. For each combination of occupation and gender, calculate the **median** age.

In [152... `df.groupby(['occupation', 'gender'])['age'].median()`

```
Out[152]:
```

| occupation | gender | |
|---------------|--------|------|
| administrator | F | 38.5 |
| | M | 35.0 |
| artist | F | 30.0 |
| | M | 32.0 |
| doctor | M | 45.0 |
| educator | F | 40.5 |
| | M | 44.0 |
| engineer | F | 29.5 |
| | M | 36.0 |
| entertainment | F | 31.0 |
| | M | 25.0 |
| executive | F | 44.0 |
| | M | 36.0 |
| healthcare | F | 43.0 |
| | M | 47.0 |
| homemaker | F | 33.5 |
| | M | 23.0 |
| lawyer | F | 39.5 |
| | M | 34.0 |
| librarian | F | 39.0 |
| | M | 38.5 |
| marketing | F | 36.5 |
| | M | 34.5 |
| none | F | 32.5 |
| | M | 16.0 |
| other | F | 34.0 |
| | M | 32.0 |
| programmer | F | 32.0 |
| | M | 30.0 |
| retired | F | 70.0 |
| | M | 61.0 |
| salesman | F | 30.0 |
| | M | 35.0 |
| scientist | F | 28.0 |
| | M | 38.0 |
| student | F | 20.0 |
| | M | 22.0 |
| technician | F | 38.0 |
| | M | 30.0 |
| writer | F | 40.0 |
| | M | 31.0 |

Name: age, dtype: float64

-
- Step 4. Create a column named 'male', which equals to 1 if gender is 'M' and zero otherwise. (tips: covert 'gender' column to numeric type: M: 1; F: 0)

```
In [153... df['male'] = df['gender']
df['male'] = df['male'].replace('M', 1)
df['male'] = df['male'].replace('F', 0)

df
```

Out[153]:

| | age | gender | occupation | zip_code | male |
|--|-----|--------|------------|----------|------|
|--|-----|--------|------------|----------|------|

| user_id | | | | | |
|---------|-----|-----|---------------|-------|-----|
| 1 | 24 | M | technician | 85711 | 1 |
| 2 | 53 | F | other | 94043 | 0 |
| 3 | 23 | M | writer | 32067 | 1 |
| 4 | 24 | M | technician | 43537 | 1 |
| 5 | 33 | F | other | 15213 | 0 |
| ... | ... | ... | ... | ... | ... |
| 939 | 26 | F | student | 33319 | 0 |
| 940 | 32 | M | administrator | 02215 | 1 |
| 941 | 20 | M | student | 97229 | 1 |
| 942 | 48 | F | librarian | 78209 | 0 |
| 943 | 22 | M | student | 77841 | 1 |

943 rows × 5 columns

```
In [154... def gender_to_numeric(x):  
    if x == 'M':  
        return 1  
    if x == 'F':  
        return 0  
  
df['male'] = df['gender'].apply(gender_to_numeric)  
  
df
```

Out[154]:

| | age | gender | occupation | zip_code | male |
|--|-----|--------|------------|----------|------|
|--|-----|--------|------------|----------|------|

| user_id | | | | | |
|---------|-----|-----|---------------|-------|-----|
| 1 | 24 | M | technician | 85711 | 1 |
| 2 | 53 | F | other | 94043 | 0 |
| 3 | 23 | M | writer | 32067 | 1 |
| 4 | 24 | M | technician | 43537 | 1 |
| 5 | 33 | F | other | 15213 | 0 |
| ... | ... | ... | ... | ... | ... |
| 939 | 26 | F | student | 33319 | 0 |
| 940 | 32 | M | administrator | 02215 | 1 |
| 941 | 20 | M | student | 97229 | 1 |
| 942 | 48 | F | librarian | 78209 | 0 |
| 943 | 22 | M | student | 77841 | 1 |

943 rows × 5 columns

```
In [155... import numpy as np  
  
df['male'] = np.where(df['gender'] == 'M', 1, 0)
```


df

Out[155]:

| | age | gender | occupation | zip_code | male |
|--|-----|--------|------------|----------|------|
|--|-----|--------|------------|----------|------|

| user_id | | | | | |
|---------|-----|-----|---------------|-------|-----|
| 1 | 24 | M | technician | 85711 | 1 |
| 2 | 53 | F | other | 94043 | 0 |
| 3 | 23 | M | writer | 32067 | 1 |
| 4 | 24 | M | technician | 43537 | 1 |
| 5 | 33 | F | other | 15213 | 0 |
| ... | ... | ... | ... | ... | ... |
| 939 | 26 | F | student | 33319 | 0 |
| 940 | 32 | M | administrator | 02215 | 1 |
| 941 | 20 | M | student | 97229 | 1 |
| 942 | 48 | F | librarian | 78209 | 0 |
| 943 | 22 | M | student | 77841 | 1 |

943 rows × 5 columns

-
- Step 5. What is the Male ratio per occupation? Sort it from the most to the least.

In [156... `df.groupby('occupation')['male'].mean().sort_values(ascending=False)`

Out[156]:

| | |
|---------------|----------|
| occupation | |
| doctor | 1.000000 |
| engineer | 0.970149 |
| technician | 0.962963 |
| retired | 0.928571 |
| programmer | 0.909091 |
| executive | 0.906250 |
| scientist | 0.903226 |
| entertainment | 0.888889 |
| lawyer | 0.833333 |
| salesman | 0.750000 |
| educator | 0.726316 |
| student | 0.693878 |
| other | 0.657143 |
| marketing | 0.615385 |
| writer | 0.577778 |
| none | 0.555556 |
| administrator | 0.544304 |
| artist | 0.535714 |
| librarian | 0.431373 |
| healthcare | 0.312500 |
| homemaker | 0.142857 |

Name: male, dtype: float64